

correlaid__text__analysis

Xiang XU, Jing(Mira) Tang, Ningze(Summer) ZU, Jianhao(Miller) Yan

November 3, 2018

Scrapping webpages

We see here we scrap three articles from correlaid.com and combine them into aone data frame.

These three articles are :

* understand-p-values; * blockchain-explained; * music-with-r.

```
## 'data.frame':   3 obs. of  2 variables:
## $ article: chr  "doc1" "doc2" "doc3"
## $ text : chr  "Skip to main content\n\nToggle navigation Menu\n\n• Zurück zu Correlaid.org\n•"
```

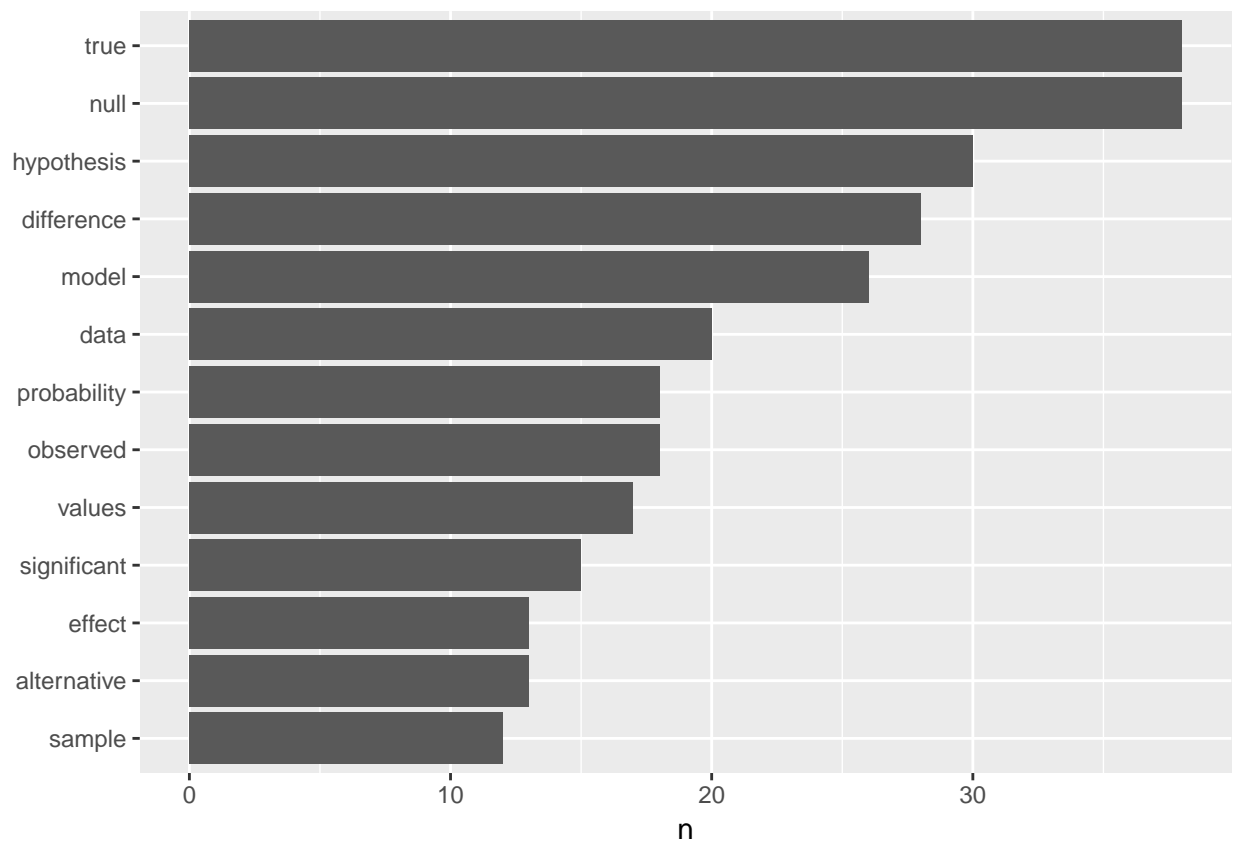
tidy text

We tidy the text by `unnest_tokens` function and filter the useful words, which make sense.

```
##   article      word
## 1   doc1      skip
## 2   doc1      main
## 3   doc1    content
## 4   doc1    toggle
## 5   doc1 navigation
## 6   doc1      menu
```

look at single word frequency and visualize

Here look at the first article p-value and cont word frequencies.

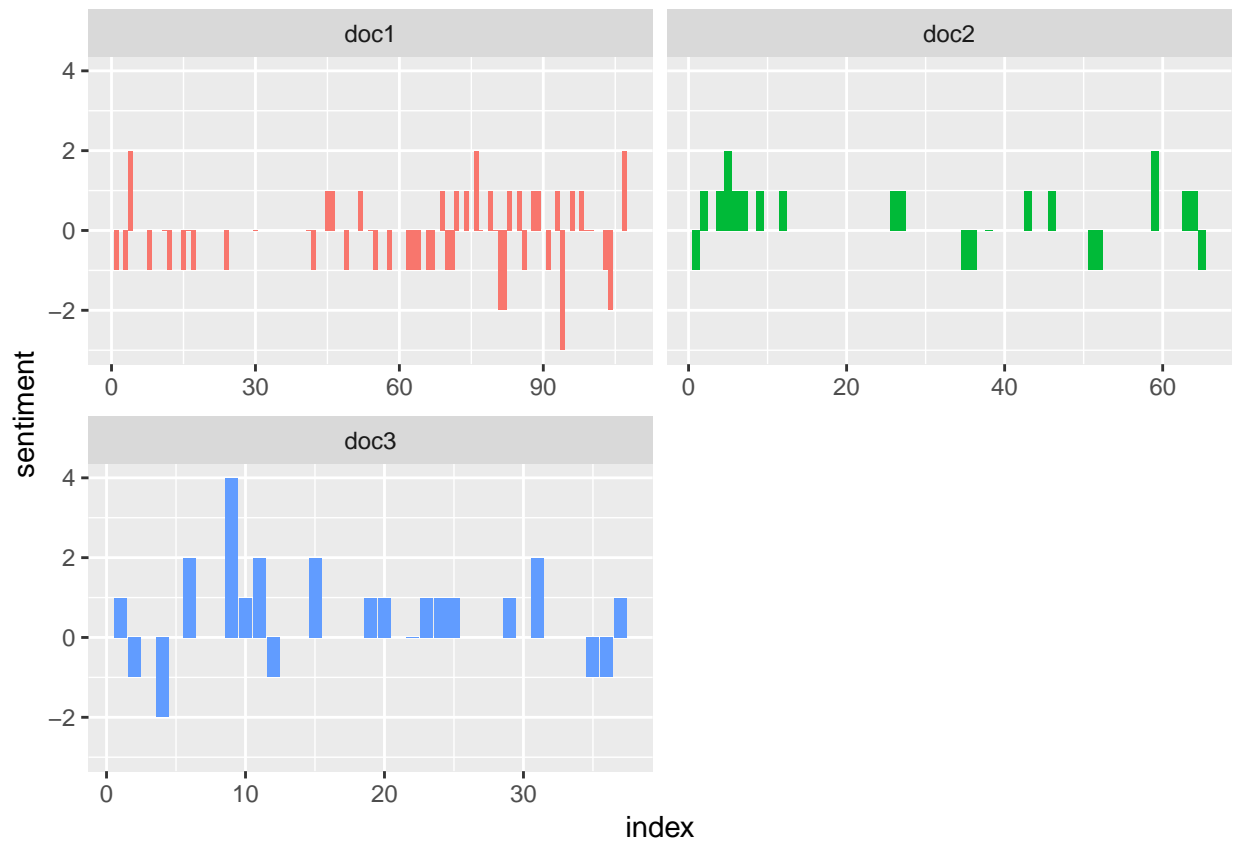


Then calculate word frequencies of three articles .

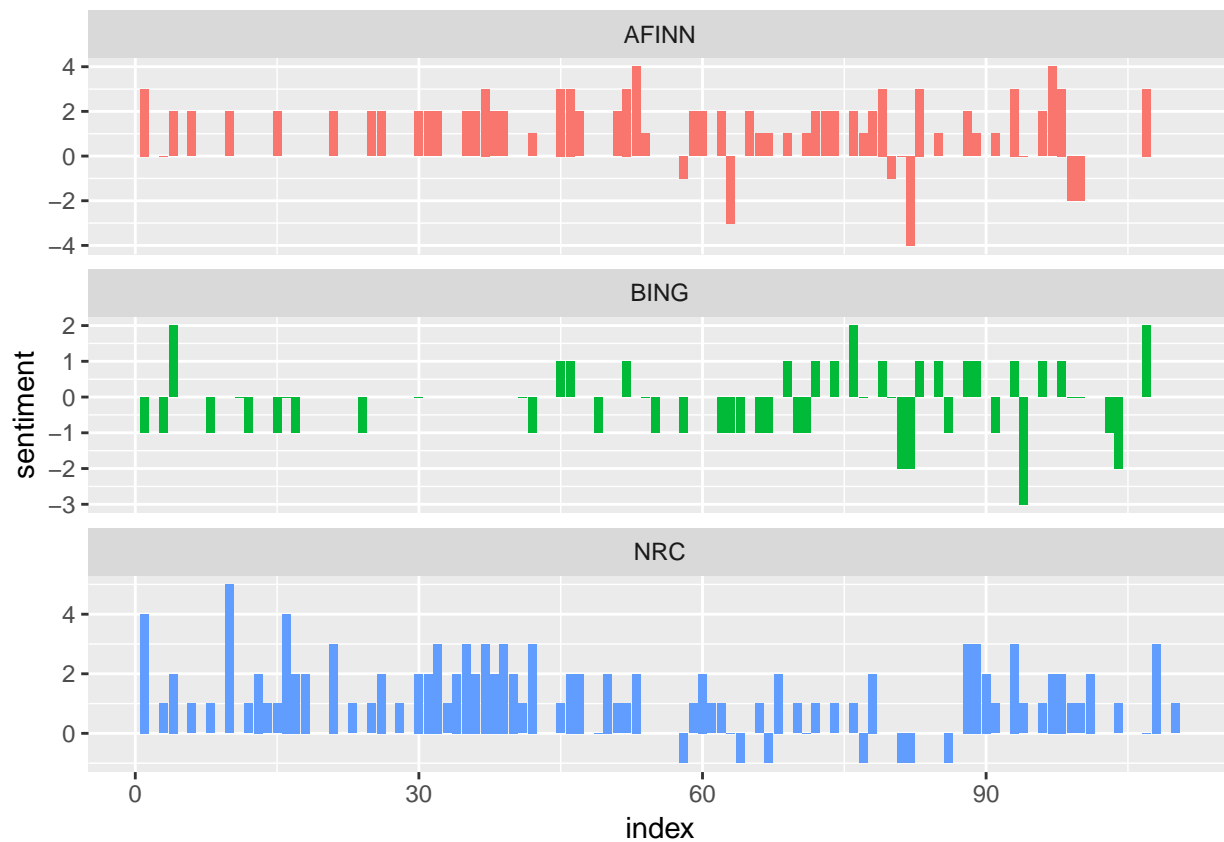
Plotting and comparing the three articles

Sentiment analysis

bing sentiment analysis

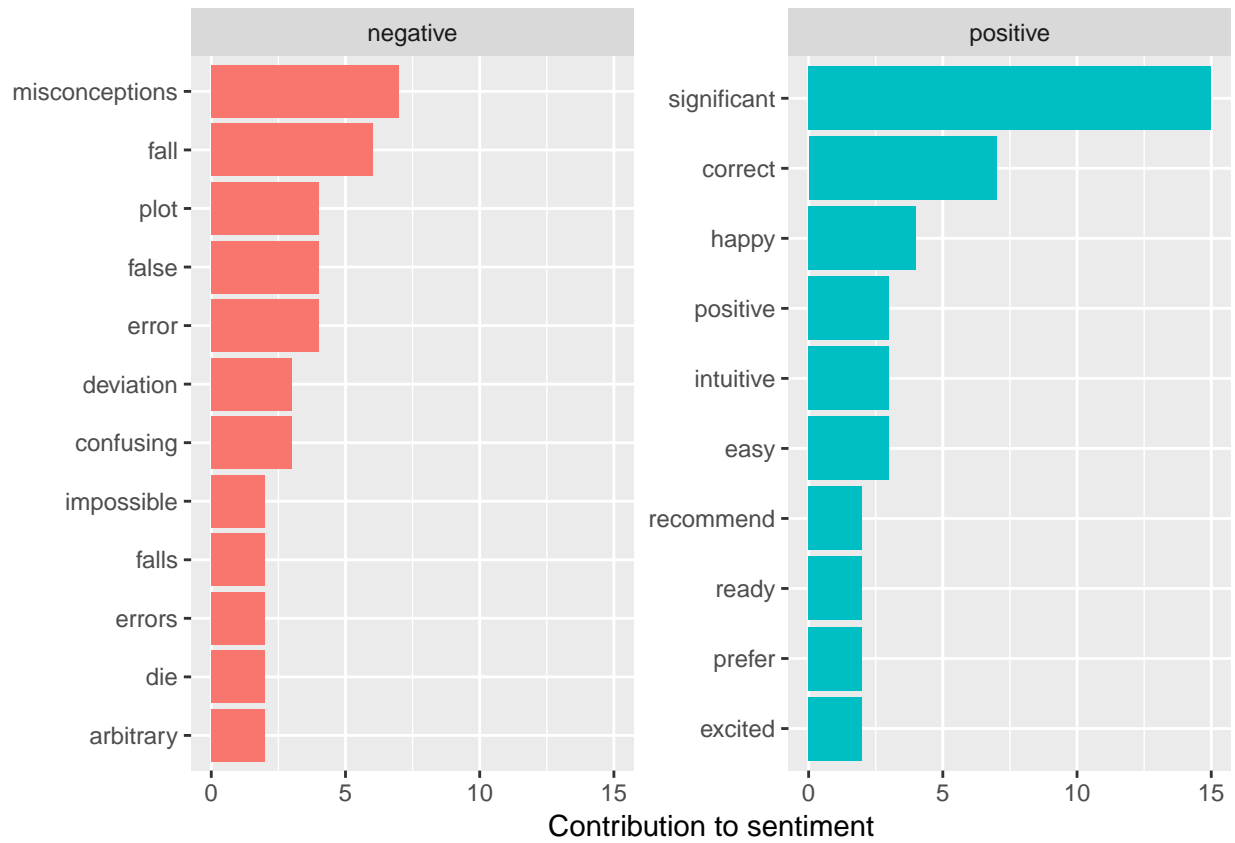


Comparing the three sentiment dictionaries



Most common sentiment words

```
## # A tibble: 79 x 3
##   word      sentiment    n
##   <chr>      <chr>    <int>
## 1 significant positive    15
## 2 correct    positive     7
## 3 misconceptions negative    7
## 4 fall        negative     6
## 5 error        negative     4
## 6 false        negative     4
## 7 happy        positive     4
## 8 plot         negative     4
## 9 confusing    negative     3
## 10 deviation    negative     3
## # ... with 69 more rows
```



Wordclouds

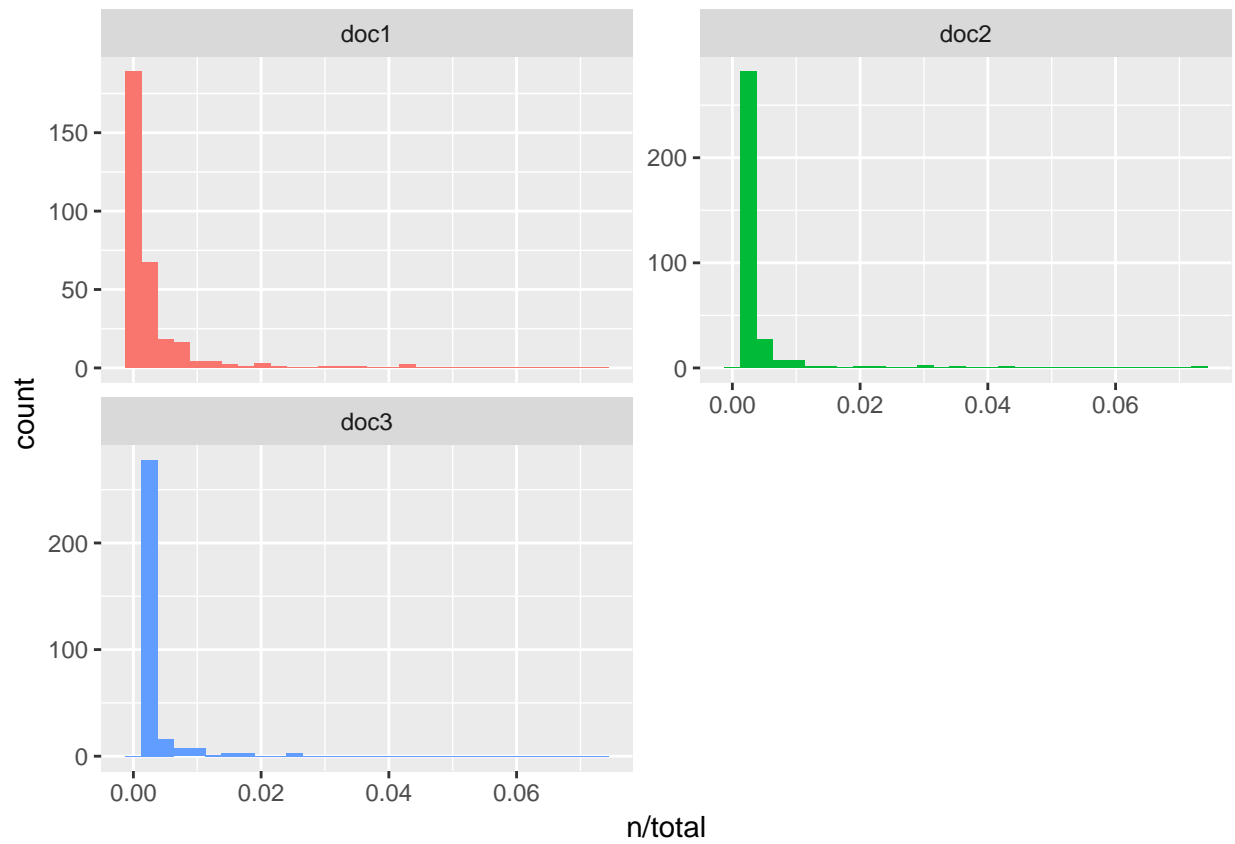


negative




positive

tf-idf

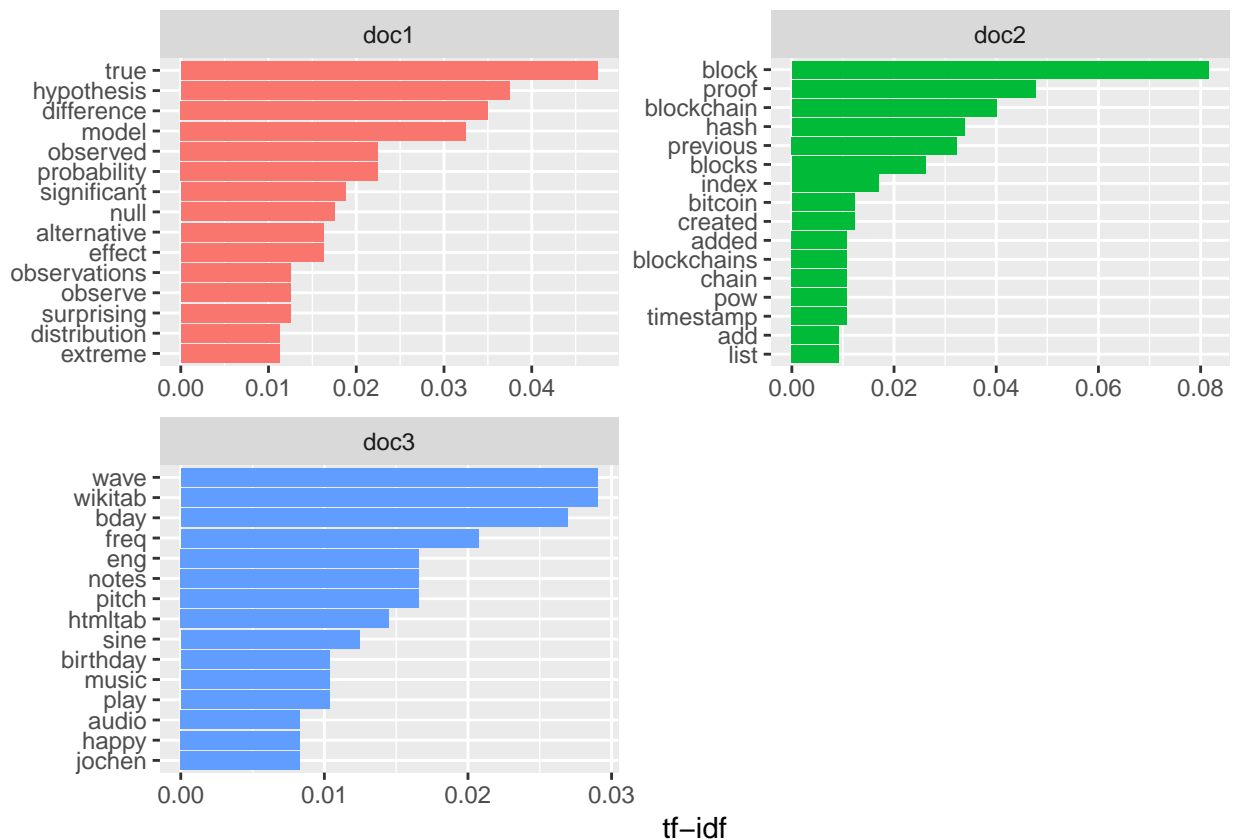


```
## (Intercept) log10(rank)
## -1.0057750 -0.8284333
```



tf-idf function

```
## # A tibble: 960 x 6
##   article word      n    tf   idf tf_idf
##   <chr>   <chr>  <int> <dbl> <dbl> <dbl>
## 1 doc2    block     53 0.0742  1.10 0.0815
## 2 doc2    proof     31 0.0434  1.10 0.0477
## 3 doc1    true      38 0.0432  1.10 0.0475
## 4 doc2    blockchain 26 0.0364  1.10 0.0400
## 5 doc1    hypothesis 30 0.0341  1.10 0.0375
## 6 doc1    difference 28 0.0319  1.10 0.0350
## 7 doc2    hash      22 0.0308  1.10 0.0339
## 8 doc1    model     26 0.0296  1.10 0.0325
## 9 doc2    previous  21 0.0294  1.10 0.0323
## 10 doc3    wave      14 0.0264  1.10 0.0290
## # ... with 950 more rows
```



n-grams and correlations

We use `unnest_tokens` function to tokenize the articles into consecutive sequences of words, called n-grams. Here we focus on bigrams, aka two consecutive words.

As one might expect, a lot of the most common bigrams are pairs of common (uninteresting) words, such as of the and to be: what we call “stop-words”. This is a useful time to use tidy’s `separate()` and `unite()`, which splits a column into multiple based on a delimiter and reunite them. In this process we can remove cases where either is a stop-word.

Also, we clean the bigrams by `str_extract()` and `filter()` function to remove cases where either is NA, space or non-letter word.

Then we look at `tf_idf` of bigrams and visualize them.

```
## # A tibble: 10 x 6
##   article bigram          n    tf   idf tf_idf
##   <chr>   <chr>      <int> <dbl> <dbl> <dbl>
## 1 doc1   null hypothesis    21 0.0843  1.10 0.0927
## 2 doc1   null model        15 0.0602  1.10 0.0662
## 3 doc2   previous block     13 0.0478  1.10 0.0525
## 4 doc1   alternative model    7 0.0281  1.10 0.0309
## 5 doc2   previous hash       7 0.0257  1.10 0.0283
## 6 doc1   omniscient jones     6 0.0241  1.10 0.0265
## 7 doc1   sample size         6 0.0241  1.10 0.0265
## 8 doc1   significant result   6 0.0241  1.10 0.0265
```

```
## 9 doc3      wikatab eng          6 0.0231  1.10 0.0254
## 10 doc1     alternative hypothesis 5 0.0201  1.10 0.0221
```

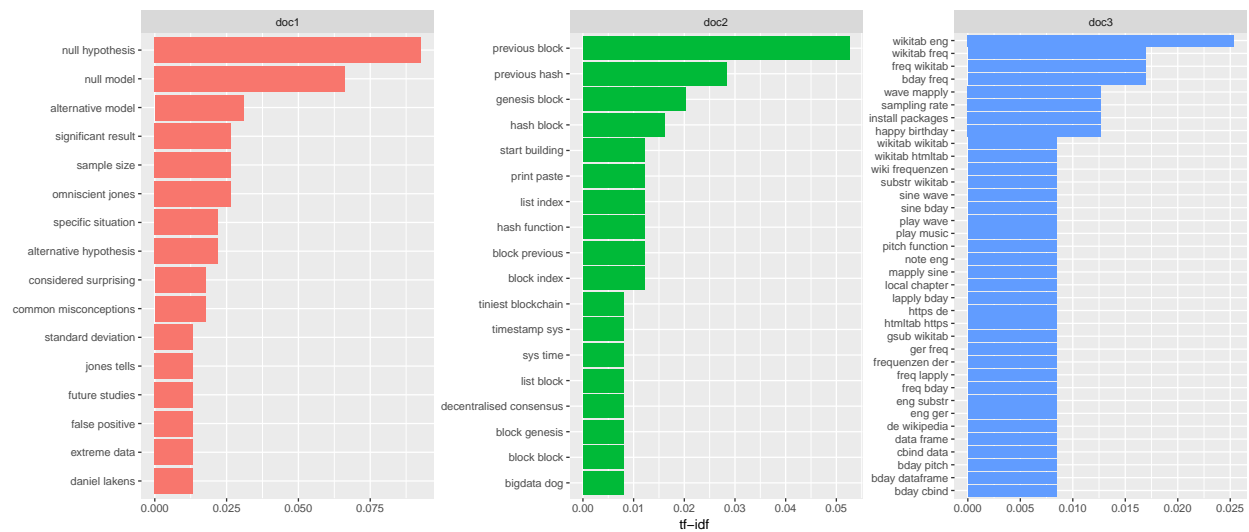


Figure: The 12 bigrams with the highest tf-idf

There are advantages and disadvantages to examining the tf-idf of bigrams rather than individual words. Pairs of consecutive words might capture structure that isn't present when one is just counting single words, and may provide context that makes tokens more understandable. However, the per-bigram counts are also sparser: a typical two-word pair is rarer than either of its component words. Thus, bigrams can be more useful when we have a larger text dataset.

- Using bigrams to provide context in sentiment analysis

```
## # A tibble: 0 x 4
## # ... with 4 variables: word1 <chr>, word2 <chr>, score <int>, nn <int>
```


For these are three academic articles and there are not many bigrams with negative terms. So we can skip this part.

- Visualizing a network of bigrams with ggraph

```
## # A tibble: 6 x 3
##   word1      word2      n
##   <chr>      <chr>    <int>
## 1 null      hypothesis    21
## 2 null      model       15
## 3 previous  block       13
## 4 alternative model       7
## 5 previous  hash        7
## 6 omniscient jones        6

## IGRAPH ffeb51b DN-- 82 62 --
## + attr: name (v/c), n (e/n)
## + edges from ffeb51b (vertex names):
## [1] null      ->hypothesis    null      ->model
## [3] previous  ->block        alternative->model
## [5] previous  ->hash        omniscient->jones
## [7] sample    ->size        significant->result
## [9] wikatab   ->eng        alternative->hypothesis
## [11] blog      ->post        genesis    ->block
```

```
## [13] specific ->situation bday ->freq
## [15] common ->misconceptions considered ->surprising
## + ... omitted several edges
```



- Counting and correlating pairs of words

```
## article word
## 1 doc1 skip
## 2 doc1 main
## 3 doc1 content
## 4 doc1 toggle
## 5 doc1 navigation
## 6 doc1 menu

## # A tibble: 294,912 x 3
## item1 item2 n
## <chr> <chr> <dbl>
## 1 main skip 3
## 2 content skip 3
## 3 toggle skip 3
## 4 navigation skip 3
## 5 menu skip 3
## 6 zur skip 3
## 7 zu skip 3
## 8 correlaid skip 3
## 9 blog skip 3
## 10 values skip 3
```

```
## # ... with 294,902 more rows
```

- Pairwise correlation

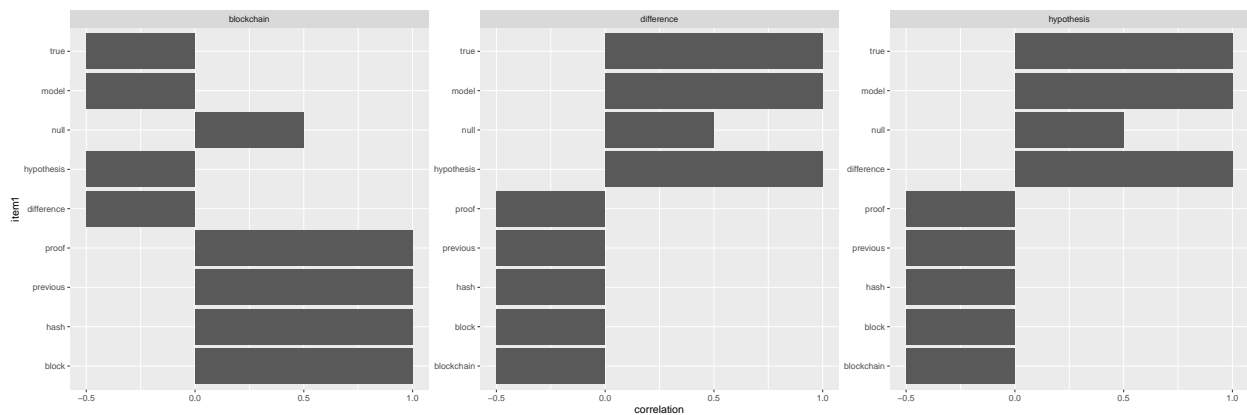
Find the phi coefficient between words based on how often they appear in the same article.

```
## # A tibble: 132 x 3
##   item1      item2      correlation
##   <chr>      <chr>      <dbl>
## 1 true      hypothesis    1.000
## 2 difference hypothesis    1.000
## 3 model      hypothesis    1.000
## 4 hypothesis true          1.000
## 5 difference true          1.000
## 6 model      true          1.000
## 7 hypothesis difference    1.000
## 8 true      difference    1.000
## 9 model      difference    1.000
## 10 hypothesis model        1.000
## # ... with 122 more rows

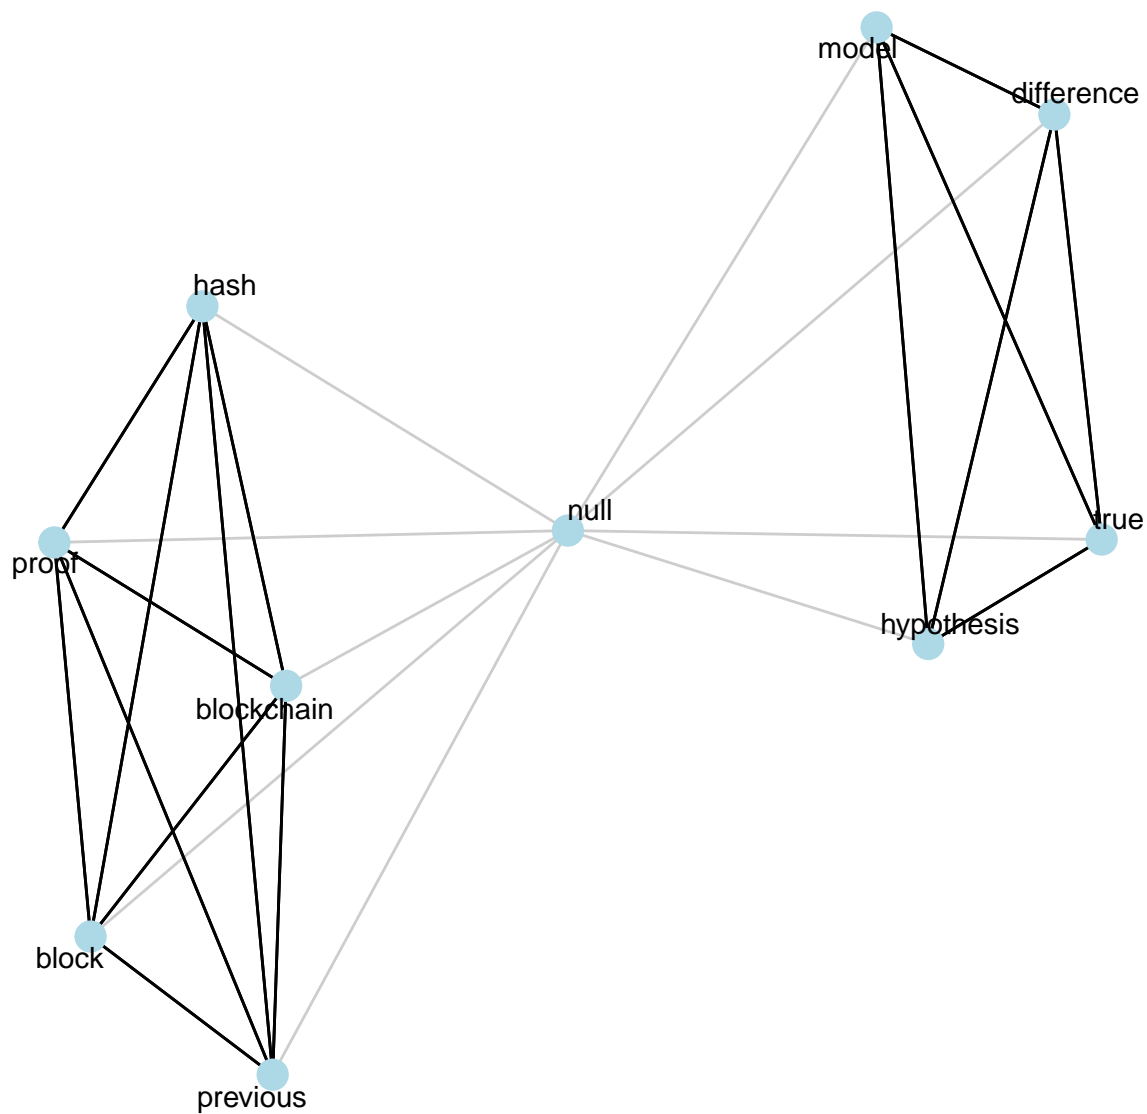
## [1] "true"      "difference" "model"      "hypothesis" "block"
## [6] "previous"  "hash"      "proof"      "blockchain" "null"
## [11] "data"      "values"

## [1] "hypothesis" "true"      "difference" "model"      "blockchain"
## [6] "block"      "previous"  "hash"      "proof"      "null"
## [11] "values"     "data"
```

Let's pick particular interesting words and find the other words most associated with them.



Visualize the correlations and clusters of words.



```
library(topicmodels)

text_dtm <- sentitext %>%
  count(article, word, sort = TRUE)%>%
  ungroup()%>%
  cast_dtm(article, word, n)

text_dtm

## <<DocumentTermMatrix (documents: 3, terms: 880)>>
## Non-/sparse entries: 1059/1581
## Sparsity          : 60%
## Maximal term length: 64
## Weighting          : term frequency (tf)

text_lda <- LDA(text_dtm,k =7, control = list(seed = 1234))
text_lda
```

```
## A LDA_VEM topic model with 7 topics.
```

```
tidy_lda <- tidy(text_lda)
tidy_lda
```


```
## # A tibble: 6,160 x 3
##   topic term      beta
##   <int> <chr>    <dbl>
## 1     1 1 block 2.66e- 2
## 2     2 2 block 1.73e-38
## 3     3 3 block 2.37e- 2
## 4     4 4 block 4.43e- 2
## 5     5 5 block 1.58e-38
## 6     6 6 block 3.97e-35
## 7     7 7 block 6.86e- 2
## 8     1 1 null  8.41e- 3
## 9     2 2 null  4.76e-36
## 10    3 3 null  2.32e- 3
## # ... with 6,150 more rows
```

```
top_terms <- tidy_lda %>%
  group_by(topic) %>%
  top_n(3, beta) %>%
  ungroup() %>%
  arrange(topic, -beta)
top_terms
```

```
## # A tibble: 22 x 3
##   topic term      beta
##   <int> <chr>    <dbl>
## 1     1 1 data    0.0798
## 2     1 1 index   0.0393
## 3     1 1 blockchain 0.0279
## 4     2 2 1      0.0414
## 5     2 2 d4     0.0245
## 6     2 2 wave    0.0215
## 7     2 2 wikipab 0.0215
## 8     3 3 1      0.0718
## 9     3 3 blocks  0.0541
## 10    3 3 data    0.0436
## # ... with 12 more rows
```

```
top_terms %>%
  mutate(term = reorder(term, beta)) %>%
  group_by(topic, term) %>%
  arrange(desc(beta)) %>%
  ungroup() %>%
  mutate(term = factor(paste(term, topic, sep = "__"),
                        levels = rev(paste(term, topic, sep = "__")))) %>%
  ggplot(aes(term, beta, fill = as.factor(topic))) +
  geom_col(show.legend = FALSE) +
  coord_flip() +
  scale_x_discrete(labels = function(x) gsub("__.$", "", x)) +
  labs(title = "Top 3 terms in each LDA topic",
       x = NULL, y = expression(beta)) +
  facet_wrap(~ topic, ncol = 4, scales = "free")
```


Top 3 terms in each LDA topic




```
lda_gamma <- tidy(text_lda, matrix = "gamma")
lda_gamma
```

```
## # A tibble: 21 x 3
##   document topic    gamma
##   <chr>    <int>    <dbl>
## 1 doc2      1 0.0446
## 2 doc1      1 0.0000584
## 3 doc3      1 0.0000873
## 4 doc2      2 0.0000705
## 5 doc1      2 0.0000584
## 6 doc3      2 0.999
## 7 doc2      3 0.231
## 8 doc1      3 0.0000584
## 9 doc3      3 0.0000873
## 10 doc2     4 0.0755
## # ... with 11 more rows
```

```
ggplot(lda_gamma, aes(gamma)) +
  geom_histogram() +
  #scale_y_log10() +
  labs(title = "Distribution of probabilities for all topics",
       y = "Number of documents", x = expression(gamma))
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Distribution of probabilities for all topics



```
ggplot(lda_gamma, aes(gamma, fill = as.factor(topic))) +
  geom_histogram(show.legend = FALSE) +
  facet_wrap(~ topic, ncol = 4) +
  scale_y_log10() +
  labs(title = "Distribution of probability for each topic",
       y = "Number of documents", x = expression(gamma))
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## Warning: Transformation introduced infinite values in continuous y-axis
## Warning: Removed 196 rows containing missing values (geom_bar).
```



Summary

By using a combination of network analysis, tf-idf, and topic modeling, we have come to a greater understanding of how datasets are related in these three articles from correlaid website. Specifically, we have more information now about how keywords are connected to each other and which datasets are likely to be related. The topic model could be used to suggest keywords based on the words in the description field, or the work on the keywords could suggest the most important combination of keywords for certain areas of study.